

Bachelor's Thesis: Performance-Analysis of VPP

Intermediate Talk

Presenter: Peter Okelmann— Advisors: Paul Emmerich, Dominik Scholz— Supervisor: Prof. Dr.-Ing. Georg Carle



VPP: a fast software router

VPP (Vector Packet Processing) is a user-space software router. This approach combines many advantages:

- ▶ deployable to usual architectures
- ▶ fast user-space network interface drivers
- ▶ can run in virtualized containers

"It is the open source version of Cisco's Vector Packet Processing (VPP) technology." [1] Now it is being developed by FD.io ("The Fast Data Project") which belongs to the Linux Foundation.

Feature Highlights:

- ▶ vectorized processing of packets in badges
- ▶ utilizes high-speed dpdk drivers
- ▶ modular and extendable packet-processing graph
- ▶ cpu-scalability

Testing Methodology

MoonGen [3] is scripted to generate testing load according to the following testing parameters:

- ▶ packet rate
- ▶ packet size
- ▶ traffic type (generic Ethernet, UDP)
- ▶ traffic pattern (inter packet gaps)

Gathered testing results:

- ▶ latency histogram
- ▶ throughput
- ▶ linux perf stats (cache misses...)
- ▶ linux perf record (cpu-time spent per symbol)
- ▶ internal vpp state information

For tests to return meaningful results, the optimum of throughput to latency is being found by a script. This packet rate is then used for further tests. Otherwise the results are inaccurate, because of utilization of the packet buffer. VPP properties to test:

- ▶ raw forwarding throughput
- ▶ latency: cache and memory impact
- ▶ processing graphs utilizing multiple processing nodes / specific processing nodes /

Measurement setup

Background

- ▶ Important background for this work
- ▶ may be some things about related work or important libraries/frameworks used

Measurement Setup:

- ▶ How does your setup look like (maybe a figure)?
- ▶ What are the relevant questions you try to answer with your measurement?
- ▶ What do you measure?
- ▶ How do you measure?

Benchmarking Setup

The diagram illustrates the benchmarking setup. It shows a 'kaunas: pos-tools' box connected to 'MoonGen' and 'Load Generator'. 'MoonGen' is connected to 'NIC1' and 'NIC2', which are connected to 'DPDK'. 'DPDK' is connected to 'I2-input' and 'I2-forward', which are connected to 'multiple cores' and 'VPP'. 'VPP' is connected to 'perf tools', which outputs 'Performance Reports'. The 'Device under Test' is also connected to the setup.

I2 Throughput

The following results are produced by sending non-ip ethernet packets to a single destination. VPP ran on an Intel E5-2640 @ 2.0GHz (cesis) with 10G networking results in the following numbers:

VPP config	max Mpps	stable Mpps	Relative
packet gen speed	14.86	14.86	100%
I2 xconnect	10.4	10.1	68%
I2 bridge: no features	9.35	9.2	62%
I2 bridge: mac-age	8.62	8.6	58%
I2 bridge: mac-learn	8.51	8.3	56%
I2 bridge: mac-learn, mac-age	8.50	8.3	56%

Planned Schedule

The graph shows latency in microseconds (ys) on the y-axis (0 to 800) versus mac flows in millions (1e7) on the x-axis (0.0 to 1.0). Multiple lines represent different percentiles of latency: 0th, 25th, 50th, 75th, 90th, 99th, and 99.9th. The 0th percentile line is flat at 0. The 99.9th percentile line shows a sharp increase in latency as mac flows increase, reaching over 800 microseconds at 1.0e7 mac flows.

Router MoonRoute FastClick (DPDK 2.2) Click (DPDK 2.2) [3] Linux 3.7

Router	Mpps	Relative
MoonRoute	14.6	100%
FastClick (DPDK 2.2)	10.4	72%
Click (DPDK 2.2) [3]	4.3	29%
Linux 3.7	1.5	10%

[1] What is vpp? https://wiki.fd.io/view/VPP/What_Is_VPP?r=1. Accessed on 2019-01-16.

[2] L. Braun, A. Didebulidze, N. Kammenhuber, and G. Carle. Comparing and Improving Current Packet Capturing Solutions based on Commodity Hardware. In *Internet Measurement Conference 2010 (IMC'10)*, Melbourne, Australia, Nov. 2010.

[3] P. Emmerich, S. Carstenmiller, D. Rauner, F. Wohlfart, and G. Carle. MoonGen: A Scriptable High-Speed Packet Generator. In *Internet Measurement Conference 2015 (IMC'15)*, Tokyo, Japan, Oct. 2015.

[4] P. Emmerich, D. Rauner, F. Wohlfart, and G. Carle. A Study of Network Stack Latency for Game Servers. In *Proceedings of the 13th Annual Workshop on Network and Systems Support for Games*, Dec. 2014.

[5] S. Günther, M. Riemensberger, and W. Utschick. Efficient GF Arithmetic for Linear Network Coding using Hardware SIMD Extensions. In *Proceedings of the International Symposium on Network Coding (NetCod)*, Aalborg, Denmark, June 2014.

Short time schedule for the upcoming weeks:

- ▶ Official start date: October 15, 2010
- ▶ Official end date: February 15, 2011
- ▶ Weeks left: 8

Schedule

- ▶ Week 1-4: Providing cookies for 8
- ▶ Week 5-6: Perform additional measurements
- ▶ Week 7: Writing thesis
- ▶ Week 8: **Several** corrections passes
- ▶ Week 9: Print and hand-in